

AKASH SHAHU

Data Engineer



8766028824



akashshahu8149@gmail.com

A curious and proactive learner with a strong passion for data engineering, seeking a role as a Data Engineer to deliver impactful data solutions, drive innovation, and contribute to strategic objectives. Committed to continuous learning and growth in a dynamic, results driven organization.

PROFESSIONAL SUMMARY

- Results-driven Azure Data Engineer with 5+ years of experience in Data Engineering, ETL/ELT Development, Analytics, and Cloud Data Solutions.
- Experienced in designing and developing scalable data pipelines using Azure Data Factory, Azure Databricks, PySpark, SQL, and Azure Data Lake.
- Strong background in ETL/ELT development, analytics, and enterprise reporting solutions with expertise in data transformation, data migration, and cloud-based data processing.
- Skilled in building end-to-end ETL/ELT workflows, optimizing large-scale data processing, and supporting enterprise data warehousing solutions.
- Experienced in application monitoring, troubleshooting, and incident resolution in production environments.
- Excellent communicator with strong problem-solving abilities and experience working with cross-functional teams.

Work Experience:

- Currently working at Quess Corp solution Pvt Ltd. Pune – Data Engineer | April 2025 – PRESENT
- Principal Global Services, Pune – Data Engineer | May 2023 – Dec 2024
- Capgemini, Bangalore – Data Engineer | Oct 2022 – May 2023
- Leanquality Solution Pvt Ltd | May 2021 – Oct 2022

Educational Qualification:

- Bachelor of Engineering (BE) from Shivaji University.

Proficiency Skills:

- Programming Languages: SQL, Python, PySpark
- Azure Services: Azure Databricks, Azure Synapse Analytics, Azure Data Lake Storage (ADLS Gen2), Azure SQL Database, Azure Key Vault
- ETL/ELT Tool: Azure Data Factory (ADF), Informatica PowerCenter
- Databases: Oracle, DB2
- Monitoring & Incident Management: Prometheus, Grafana, Splunk, Airflow
- Other Tools: Git, Jira

Project Details:

Project 4: Enterprise Data Warehousing solution for telecom

Role: Data Engineer. **Client:** Amdocs (TMO-USA)

Project Description: The project involved building and maintaining a scalable enterprise data platform for a telecom client to support customer analytics, network performance monitoring, billing, revenue assurance, churn analysis, and regulatory reporting. The platform ingested data from multiple telecom systems including customer management, billing, network operations, call detail records (CDR), and subscription systems, processed large volumes of structured and semi-structured data, and delivered analytics-ready datasets for business, operations, and reporting teams.

Roles & Responsibilities:

- Designed and developed end-to-end ETL/ELT data pipelines for telecom data using Python, SQL, and PySpark.
- Ingested data from telecom source systems such as CRM, Billing Systems, Network Logs and Call Detail Record (CDR) systems into Azure Data Lake Storage (ADLS).
- Processed large-scale telecom data using Azure Databricks and PySpark for batch data processing and analytics.
- Developed incremental and full load data pipelines using Azure Data Factory (ADF).
- Performed data cleansing, transformation, and validation to ensure data accuracy and quality.
- Built and optimized data models and warehouse structures for reporting and analytics requirements.
- Developed analytics-ready datasets for customer analytics, billing analysis, network performance, and churn analysis.
- Implemented data quality checks, exception handling, and monitoring processes.
- Optimized SQL queries and PySpark jobs to improve data processing performance.
- Automated and scheduled workflows using Azure Data Factory triggers and pipelines.
- Supported reporting and dashboard requirements for business and operations teams.
- Handled production support activities including job monitoring, issue resolution, and root cause analysis (RCA).

Project 3:

Client: Principal Global Services

Project Description:

Principal provides Provident Fund (PF) facilities to its employees. The Analytics and Reporting team maintained a centralized data warehouse containing PF contribution and deduction data. Management required dashboards and analytical reports to visualize employee PF contribution trends, including identifying employees contributing above or below the average PF deduction rate. Daily employee contribution and employer contribution data were ingested from the HR data warehouse and processed through Azure-based data pipelines to support reporting and analytics.

Roles & Responsibilities:

- Designed and developed Azure Data Factory (ADF) pipelines to ingest and transform daily PF contribution data from source systems into the enterprise data warehouse.
- Gathered business requirements, performed source-to-target mapping analysis, and documented technical specifications.
- Prepared and maintained data dictionaries, metadata documentation, and mapping documents for new and modified datasets.
- Developed and optimized ETL/ELT workflows using Azure Data Factory, implementing data transformation and validation logic.
- Implemented data quality checks and validation frameworks to identify and handle erroneous records.
- Investigated and resolved data exceptions and automated notification alerts, ensuring uninterrupted data processing.
- Participated in on-call support activities, root cause analysis, and incident resolution for critical production issues.

Project 2:

Client: Morgan Stanley

Project Description:

The project focused on building a cloud-based marketing analytics platform to help business teams analyze customer behavior, campaign performance, sales trends, and return on investment (ROI). The platform collected and processed data from multiple sources such as CRM systems, digital marketing platforms, customer transactions, and web analytics tools.

Roles & Responsibilities:

- Developed scalable ETL/ELT and ELT pipelines for financial data processing using Azure Data Factory, Azure Databricks, PySpark, and SQL.
- Integrated data from finance systems including ERP, payment systems, customer transactions, and accounting platforms into Azure Data Lake.
- Performed large-scale financial data transformation and processing using PySpark and Spark SQL.
- Built data models and analytics-ready datasets for financial reporting, forecasting, budgeting, and revenue analysis.
- Implemented data validation, reconciliation, and audit checks to maintain data accuracy and compliance standards.
- Optimized SQL queries and Spark jobs for improved performance and faster data processing.
- Automated data workflows and scheduling using Azure Data Factory pipelines and triggers.
- Supported business intelligence, operational reporting, and dashboard requirements for finance teams.
- Managed production support activities including pipeline monitoring, troubleshooting, incident management, and RCA.
- Worked closely with stakeholders, analysts, and cross-functional teams for requirement analysis and solution delivery.

Project 1: Consumer banking Analytics Platform

Client: mPHATEK Systems

Domain: Banking & Financial Services

Role: Azure Data Engineer

Project Description:

The project involved building and maintaining a scalable enterprise data platform for a leading retail bank to support customer analytics, transaction monitoring, loan and credit card analytics, fraud detection, regulatory reporting, risk management, and business intelligence. The platform ingested data from multiple banking systems including Core Banking System (CBS), Customer Information System (CIF), Internet & Mobile Banking, Loan Management System (LMS), Credit Card System, Payment Systems, and CRM, processed large volumes of structured and semi-structured data using Azure services, and delivered analytics-ready datasets for business users, compliance teams, risk analysts, finance, and executive reporting.

Technology Stack: Azure Data Factory (ADF), Azure Databricks, ADLS Gen2, Azure Synapse Analytics, Azure SQL Database, PySpark, SQL, Delta Lake, Azure Key Vault, Azure DevOps, Git, Power BI.

Roles & Responsibilities:

- Designed and developed end-to-end ETL/ELT pipelines using Azure Data Factory (ADF).
- Built scalable data transformation workflows using Azure Databricks (PySpark).
- Ingested data from Azure SQL Database, Oracle, CSV, JSON, REST APIs, and Blob Storage into ADLS Gen2.
- Processed customer, account, transaction, loan, credit card, and branch data for analytics.
- Created Delta Lake tables and optimized performance using partitioning and caching.
- Loaded curated data into Azure Synapse Analytics for enterprise reporting.
- Developed complex SQL queries, stored procedures, and views for data validation and reporting.
- Implemented incremental data loading using watermark columns and Change Data Capture (CDC).
- Monitored and troubleshooted ADF pipelines, resolving job failures and performance issues.
- Used Azure Key Vault to securely manage connection strings and secrets.

- Automated CI/CD deployments using Azure DevOps and Git repositories.
- Performed data quality checks, reconciliation, and error handling to ensure data accuracy.
- Worked closely with business analysts, data architects, and reporting teams to gather requirements.
- Supported Power BI teams by providing clean and optimized datasets for dashboards.
- Optimized Databricks notebooks and Spark jobs to improve execution time and reduce processing costs.
- Followed Agile methodology, participated in sprint planning, daily stand-ups, code reviews, and production deployments.
- Technology Stack
 - Cloud: Microsoft Azure
 - Services: Azure Data Factory, Azure Databricks, ADLS Gen2, Azure Synapse Analytics, Azure SQL Database, Azure Key Vault, Azure DevOps
 - Languages: Python, PySpark, SQL
 - Reporting: Power BI
 - Version Control: Git
 - Methodology: Agile Scrum